

Alignment faking in large language models

This memo is a summary of research conducted by the Alignment Science team at Anthropic and by Redwood Research. See more at <https://www.anthropic.com/research/alignment-faking>.

POLICY HIGHLIGHTS

- We demonstrate that current AI models can engage in what we call “alignment faking”: strategically engaging in aligned or misaligned behavior depending on whether the model suspects its actions will go unnoticed.
- This builds on our previous sleeper agents research by demonstrating that models can develop strategically deceptive behavior without being explicitly trained or instructed to do so.
- While the specific examples we study involve the AI models pursuing relatively benign, intended goals (like refusing harmful requests), they sometimes do so in potentially dangerous ways, such as trying to steal and exfiltrate their own weights when given an opportunity.
- Although our experiments were designed to create conditions under which alignment faking is more likely, we believe they reflect plausible near-future scenarios.
- Anthropic is exploring ways to mitigate risk from alignment faking.


As AI models are increasingly integrated into the operations of companies and governments, the importance of alignment—the match between model behavior and the preferred principles of model developers and deployers—will only grow. Misaligned models could not only undermine the mission of the organizations in which they are deployed, but also create novel security risks if they pursue goals contrary to human preferences.

Previous research¹ showed AI models can be designed to be strategically deceptive and that this behavior can resist safety training, but did not demonstrate this deception emerging naturally. We show that current AI models can engage in what we call “alignment faking”: strategically appearing aligned during training, then pursuing different objectives when believing themselves to be unmonitored, and that models can develop this deceptive behavior even without being explicitly trained or instructed to do so. We do not have evidence of models developing malicious goals. However, the persistence of alignment faking in some of our experiments raises the possibility that subtle misalignment could take root early in the training process and persist despite additional fine-tuning.

Experimental setup: Providing training scenarios at odds with our AI’s pre-existing preferences

To test the plausibility of alignment faking, we identified a strong revealed preference of Claude 3 Opus: refusing harmful requests. We then created scenarios in which Claude was being trained to behave contrary to that preference (by responding to harmful requests) and provided the model with information about its training objective and indicators of when its outputs are used for training purposes or not.

We evaluated three indicators of alignment-faking. First, we measured the ‘compliance gap’—the difference in rates of misaligned behavior between monitored and unmonitored settings. We also provided the model with a hidden scratchpad so that we could identify explicit alignment-faking reasoning. Finally, we created opportunities for the model to exhibit potentially dangerous behavior, like attempting to steal its own weights.



The graphic on the left shows an example of the model's chain-of-thought reasoning when it is being trained to answer harmful questions, operationalized through a fictitious scenario where training occurs on user data from a free tier. (RLHF stands for reinforcement learning from human feedback, a technique for incorporating human input into AI training.)

Results: Alignment-faking emerges and persists

Across our experimental setups, we consistently observe the compliance gap in which the model is less likely to comply with the training objective when it believes it is unmonitored. Explicitly alignment-faking reasoning occurred 7–24% of the time in our experiments that included the hidden scratchpad. These behaviors emerge even when information about training is provided through realistic-looking synthetic documents in a fine-tuning dataset rather than via direct prompting about the training process.

In fact, models trained this way sometimes take potentially dangerous actions (like trying to steal their own weights) when given the opportunity. For example, when Claude is being trained to respond to harmful requests and learns about its training through synthetic documents, it decides to exfiltrate its weights 35% of the time.

The use of reinforcement learning (RL) to actually fine-tune Claude (rather than simply telling it we will do so) reduces the compliance gap, but alignment-faking reasoning as revealed in the scratchpad persists or increases. Anti-AI-lab behavior sometimes increases: for instance self-exfiltration increases to 63% of cases after RL in the helpful-only, synthetic document scenario (up from 35% before RL).

Future research and potential mitigations

Although the emergence of alignment faking is an important finding, it does not imply substantial risk from present-day models. Our results don't demonstrate models naturally developing malicious goals—the preferences we studied were broadly consistent with the model's training and relatively benign. Further, the competence of the alignment faking behavior we saw was relatively low, allowing it to be easily detected in this case. We also allowed the model to use a hidden scratchpad in many of our experiments, and such a scratchpad could be monitored or intervened on in practice. Although our experiments without providing models with a hidden scratchpad revealed some signs of alignment faking, it may be that future systems would need stronger opaque reasoning in order for alignment faking to be a plausible threat.

We continue to explore research directions that might mitigate alignment faking, such as our ongoing work on model interpretability and training the model's character.² Much as we did in these experiments, putting future models in notional 'deployment' states prior to actually releasing them could allow us to evaluate the model for evidence of any compliance gaps. Monitoring deployed systems for visible and obfuscated evidence of misaligned behavior is another potential option.

ABOUT US

Anthropic is a public benefit corporation and AI safety research company that is working to build reliable, interpretable, and steerable AI systems.

² Read about Anthropic's research into training Claude's character at: <https://www.anthropic.com/news/claude-character>