

The Enterprise AI Transformation Guide

Accelerate your enterprise AI transformation with proven strategies from Anthropic's customers and internal teams.



Contents

Foreword	3
Step 1: Lay the foundation	5
Step 2: Launch a pilot	9
Step 3: Scale impact	12
Building an AI-first enterprise: real-world examples from Anthropic	15
Your AI transformation starts today	18

The AI imperative: why the time to act is now.

92% of companies plan to invest in generative AI (GenAI) over the next 3 years, yet only 1 percent believe their investments have reached full maturity.

But investment alone won't get you there.

The good news? You're not alone. Working alongside industry leaders like Cox Automotive, Thomson Reuters, and NBIM, Anthropic has developed a proven three-step approach to driving enterprise AI transformation at scale:

- **Lay the foundation:** Organizations begin by building a solid AI foundation that includes strategy development, stakeholder alignment, and governance.-

- **Launch a pilot:** Next, enterprises launch carefully selected pilots that demonstrate value quickly while building expertise.

- **Scale impact:** Finally, successful pilots scale across the business through structured rollout programs and centers of excellence.

In this guide, we'll walk through how to implement these three steps and in the process:

- Identify high-impact use cases specific to your industry
- Build effective business cases that secure cross-functional support
- Drive AI upskilling and adoption across your organization

We'll also share how teams at Anthropic deploy AI across functions, from engineering and legal teams to marketing and HR.

Let's dive in.



Chapter 1

Step 1: Lay the foundation

Step 1: Lay the foundation

Driving enterprise-wide AI transformation starts with establishing the organizational groundwork that enables sustainable adoption and measurable impact. In this section, we discuss how to define your AI strategy, drive alignment, establish a clear AI governance framework, and more best practices gleaned from our own experiences working with leaders to chart the course for their organizations' AI journeys.

Setting your AI strategy

Successful AI transformation begins with a few clear pilot projects aligned to business objectives. Identify where AI delivers greatest impact and justify the investment required.

Consider these key – and popular – functional areas that teams at Rakuten, Intuit, Thomson Reuters, and Amazon have chosen as their pilot projects when mapping high-impact use cases:

- **Engineering teams** undergo fundamental changes in software development lifecycles as agentic coding solutions like Claude Code assist with code generation, reviews, and debugging, as well as more robust tasks, such as feature prototyping and code migrations.
- **Legal departments** automate contract review, draft standard documents, and synthesize research from multiple sources.
- **Marketing organizations** leverage AI for content creation, campaign optimization, and cross-channel performance analysis.

- **Finance functions** streamline analysis, automate reporting, and improve forecasting accuracy with agentic tooling, like Claude for Financial Services.
- **Customer support teams** deploy intelligent chatbots while building AI-enhanced product features.

Each function presents unique opportunities based on your company's current pain points and potential for sustained improvement. The key is identifying where AI can deliver meaningful improvement rather than attempting to transform everything simultaneously.

Driving leadership and stakeholder alignment

Executive alignment is critical for any type of enterprise transformation, particularly when it comes to adopting a technology as powerful as AI. Leaders must understand both the opportunities AI presents—including the potential for 20-30% gains in productivity, speed to market, and revenue according to PwC—alongside challenges such as substantial upfront investments, cultural resistance, and organizational change barriers. This alignment extends beyond initial approval to ongoing support through inevitable implementation challenges.

Technical solutions alone cannot drive transformation—people and processes must evolve alongside these implementations. Organizations that treat change management as an afterthought often struggle with adoption.

Assembling your AI steering committee

The key to successful change management begins with assembling a steering committee that represents critical business functions and decision-making authority. This group should include the CEO or another C-suite sponsor who can remove organizational obstacles, functional leaders who understand

operational realities, technology executives who grasp implementation requirements, finance representatives who track ROI and manage budgets, and legal or compliance leaders who can establish governance frameworks.

Building AI implementation champions

In addition to assembling a larger steering committee, it's important to identify and empower change champions at every organizational level. These individuals should include respected managers who influence their peers, technical experts who understand both legacy systems and AI capabilities, early adopters enthusiastic about innovation, and skeptics whose questions reveal legitimate implementation concerns. Provide champions with additional training, direct access to leadership, and recognition that elevates their status while making their advocacy visible across the organization.

Customer spotlight: Quantium, a global data and AI technology solutions provider, uses Claude to drive company-wide AI adoption, with over 89% using AI in their daily work.

Aligning on clear success metrics & timelines

Before launching any pilot, establish concrete success metrics that stakeholders understand and accept. These metrics typically span four dimensions:

- **Adoption metrics** track how many users engage with AI tools and how frequently they use them—for example, measuring daily active users, feature utilization rates, and session frequency across different departments.
- **Efficiency measures** document time saved and productivity improvements, such as reducing contract review time from 2 hours to 30 minutes or enabling customer service agents to handle 13.8% more inquiries per hour.

- **Quality metrics** ensure AI outputs meet organizational standards through measures like 95% accuracy in data extraction, error rates below 5%, or first-pass approval rates for AI-generated content.
- **Satisfaction scores** capture user experience and identify areas needing improvement via Net Promoter Scores (NPS), task difficulty ratings, and willingness to recommend the tool to colleagues.

Most successful pilots run 8-12 weeks—long enough to see meaningful patterns, short enough to maintain momentum. Expect the first 2-3 weeks for onboarding and adjustment as users learn the tool. Measurable efficiency gains typically emerge by week 4-6. Quality improvements and adoption patterns become clear by week 8-10. Resist the urge to extend pilots indefinitely; if you're not seeing results by week 12, you likely need to adjust your approach or use case, not your timeline.

Track these metrics weekly to catch issues early, review monthly to identify trends, and adjust your approach based on data rather than assumptions. This creates accountability while building confidence in your AI initiatives.

Customer spotlight: Zapier, the leading AI orchestration platform, demonstrated metrics that balance employee adoption with value, tracking 800+ AI agents deployed and 10x year-over-year growth in Claude usage.

Establish measurement discipline

To make it easier to measure the pilot's impact and adoption, implement automated data collection wherever possible to ensure measurement accuracy and reduce the burden of manual reporting (hint: this is a great use case for AI). Deploy usage analytics that track every AI interaction—time spent, features accessed, tasks completed, and error rates encountered. Consider establishing a pilot review cadence, including:

- **Regular stakeholder reviews** that translate raw metrics into actionable insights. Weekly dashboards should highlight adoption trends like showing customer service adoption grew week-over-week, identify struggling user segments, and flag any quality issues.
- **Monthly business reviews** should present trend analysis showing whether pilot trajectories will meet success thresholds, with early warning indicators like declining user engagement after the initial month or plateauing efficiency gains.
- **Quarterly executive briefings** should synthesize findings into strategic recommendations. Include qualitative feedback alongside quantitative metrics like user testimonials and specific use cases.

Customer spotlight: IG Group, the global online trading provider, adopted a manager-led AI adoption program with a tiered training program and specific implementation targets, leading to 60%+ adoption of Claude in the first three months.

Establish an AI governance framework

AI governance frameworks are a critical component of any enterprise AI transformation by enabling innovation while managing risk through carefully designed policies that balance protection with productivity. Here are some critical components:

- **Access controls** determine who can use AI systems and what data they can access, implementing role-based permissions that align with job responsibilities and information sensitivity levels.
- **Usage guidelines** clarify acceptable applications while explicitly prohibiting problematic uses such as processing personally identifiable information in public AI models, making employment decisions without human oversight, or generating content that could expose the organization to legal liability.

- **Quality standards** specify review requirements and accuracy thresholds, defining when AI outputs require human verification versus when they can proceed autonomously.
- **Compliance protocols** ensure AI use meets regulatory requirements across all relevant jurisdictions, from GDPR in Europe to industry-specific regulations in healthcare and financial services.

The earlier you prioritize governance and compliance, the more successful your AI transformation will be in the long-term.

The importance of AI governance: AI governance is core to our DNA; in fact, we were the first AI company to achieve the ISO 42001 for responsible AI. Many of our customers have found the following resources useful when it comes to better understanding the foundations of AI safety and governance:

Claude's Constitution: Details the principles that guide Claude's behavior through our Constitutional AI approach

Constitutional AI: Research paper that highlights Anthropic's foundation for training AI systems with embedded values and principles.

Anthropic's Responsible Scaling Policy: Discusses Anthropic's responsible scaling policy, a risk governance framework we use to mitigate potential catastrophic risks from frontier AI systems.



Chapter 2

Step 2: Launch a pilot

Step 2: Launch a pilot

Successful pilots deliver quick wins while building organizational capability. Choose projects carefully, showcase cross-functional potential, and learn rigorously from every experiment.

Choose initial projects carefully

We suggest selecting one or two initial projects that demonstrate AI's value while building organizational capability. These pilots should span different business functions to showcase AI's versatility. For example, building an internal chatbot might answer common IT and support questions, reduce help desk burden, or deploying a coding agent like Claude Code could accelerate software development while improving code quality across multiple business units.

Pro-tip: Demonstrate value quickly through sprint-based implementation. Most successful pilots show meaningful results within 30-60 days rather than requiring months of development. This rapid value demonstration builds momentum and justifies continued investment.

When choosing your pilot project, ensure they abide by these two critical guidelines:

- **Evaluate potential pilots against clear ROI:** Measure the potential business impact in specific terms like cost reduction or time savings, technical feasibility, user receptivity or process disruption, and resource requirements including both technology costs and implementation effort.

- **Prioritize projects where failure creates minimal business disruption:** avoid selecting customer-facing applications or mission-critical processes for initial pilots regardless of their potential value.

Once you've aligned on your chosen pilot, structure pilot teams with clear roles and dedicated time commitments rather than treating AI as an additional responsibility for already-stretched resources. Assign a pilot lead who owns outcomes and coordinates across functions, technical resources who handle implementation and troubleshooting, business users who provide domain expertise and test real-world scenarios, and executive sponsors who remove organizational barriers and maintain stakeholder alignment.

Showcase cross-functional potential

Design pilots that reveal possibilities beyond their immediate scope. When legal teams see marketing's success with content generation, they begin imagining contract automation applications. When engineering demonstrates code review acceleration, finance envisions automated report generation. This cross-pollination of ideas drives organic expansion beyond initial implementation plans.

Create structured opportunities for cross-functional learning that accelerate insight transfer across the organization, such as a monthly "AI Showcase" where pilot teams highlight their projects. These sessions should include live demonstrations of AI capabilities, before-and-after comparisons of work processes, and open discussion about challenges encountered and solutions discovered.

Customer spotlight: GitLab, a leading DevSecOps platform, used Claude for Work to drive productivity across engineering, sales, and marketing teams, creating forums for employees to share best practices and learnings.

Customer spotlight: Bridgewater Associates, one of the world's largest hedge funds, deployed Claude for Work across research and operations teams. When adoption lagged in certain departments, they created peer mentorship programs where successful users helped colleagues discover relevant use cases, dramatically accelerating organization-wide engagement.

Navigate common pilot challenges

Even well-designed pilots encounter obstacles. Anticipating common challenges and preparing response strategies prevents minor issues from derailing valuable initiatives. We outline several of the most prevalent, below.

- **User resistance and adoption gaps** emerge when team members feel threatened by AI or find new workflows disruptive. Address this through early involvement in pilot design, hands-on training that builds confidence rather than just explaining features, and celebrating early adopters who become internal advocates. When a team member discovers a better way to use the technology, amplify their insight across the organization.
- **Data quality and availability issues** often surface mid-pilot when teams discover their existing data isn't structured for AI applications. Rather than pausing the pilot, implement pragmatic workarounds such as starting with better-quality data subsets while simultaneously improving broader data practices, or using AI to help clean and structure the problematic data itself.
- **Integration complexity** with legacy systems frequently exceeds initial estimates. Build in technical flexibility by starting with manual handoffs between AI and existing systems if necessary, then automating connections as you prove value. Perfect integration shouldn't block initial learning.
- **Scope creep** threatens pilot timelines when teams discover additional applications mid-implementation. Maintain discipline by documenting new ideas in a "next phase" backlog rather than expanding current scope. This preserves pilot timeline integrity while ensuring good ideas aren't lost.

Conduct pilot post-mortems

Once each pilot concludes, the real work of learning begins, moving beyond surface metrics to understand the deeper story of what unfolded and why.

Start by examining the numbers but also evaluate quantitative metrics. In fact, these “anecdotal” proof points are often more informative, for instance, moments when users discovered unexpected benefits, the friction points that emerged in practice, the workarounds teams invented when the technology didn't quite fit their workflow.

On the technical side, probe system reliability, integration snags, data quality surprises, and infrastructure needs that only became apparent under real-world conditions. Understanding user adoption requires detective work—why did some teams embrace the technology while others quietly resisted? What practical barriers emerged?



Chapter 3

Step 3: Scale impact

Step 3: Scale impact

Moving from successful pilots to enterprise-wide transformation requires structured training programs, dedicated centers of excellence, and governance that scales with adoption.

Transform pilots into a launchpad for AI upskilling

Scaling success demands more than replicating what worked—it requires building genuine capability across every level of your organization. Think of training not as checkbox compliance but as a strategic opportunity to fundamentally upskill your workforce for an AI-augmented future.

Different audiences need dramatically different learning journeys, for example:

- **Your executives** don't need to master prompt engineering—they need strategic context to make sound investment decisions and understand competitive implications.
- **Managers** occupy the critical middle ground, translating strategy into daily practice.
- **Power users**—your AI champions—deserve deep technical training covering advanced features, prompt engineering nuances, and sophisticated troubleshooting.

Don't overlook the power of experiential learning. Hackathons inject energy and experimentation into what could otherwise feel like mandatory corporate training. When teams compete to solve real business problems using AI, learning happens organically while building genuine enthusiasm that top-down mandates can never achieve.

You can also create certification programs that validate competency while fueling ambition for continued growth. Design multi-level credentials from foundational literacy through advanced practitioner status, but insist on demonstrated proficiency through real-world assessments rather than multiple-choice tests that prove nothing.

Pro-tip: Recognition matters: celebrate certified users in internal communications, provide them priority support access, and most importantly, consider certification status in promotion decisions. Nothing signals organizational commitment like tying advancement to AI proficiency. This transforms training from "nice to have" into "essential for career growth"—exactly the mindset shift you need to build true institutional capability that outlasts any individual pilot's success.

Establish centers of excellence

Create specialized teams – or centers of excellence – dedicated to sustaining and expanding AI capabilities. These centers of excellence develop best practices for implementation across functions, ensuring consistent approaches and knowledge sharing, provide technical support and troubleshooting when users encounter challenges, and systematically experiment with new use cases.

Structure centers of excellence with clear accountability and cross-functional representation that prevents siloed thinking. Include technical architects who understand system integration and data flows, domain experts from each major function who translate business needs into AI opportunities,

and data scientists who optimize model performance and identify emerging capabilities.

Pro-tip: To facilitate additional upskilling, establish regular rotation programs that bring functional experts into the center for 3-6 month increments, building their AI expertise while ensuring the center stays connected to evolving business needs.

Measure and amplify success across the organization

Pilot wins mean nothing if they stay isolated. Transform individual successes into enterprise momentum by systematically measuring impact and broadcasting results that inspire broader adoption.

Develop a comprehensive ROI framework that captures value beyond simple cost savings. Financial impact matters—track hard savings from reduced headcount needs, vendor costs, or operational expenses, plus productivity gains measured in time saved and output increased. But equally important are strategic benefits like faster time-to-market for new products, improved decision quality from better data analysis, and enhanced employee satisfaction from eliminating tedious work.

Here are a few other ways to communicate these success stories:

Create compelling narratives around quantified results.

Numbers alone rarely inspire action—executives and employees need stories that bring data to life. When your sales team reduces proposal turnaround from 5 days to 2 hours using AI, don't just report the metric—share the story.

Structure regular "wins reporting" that maintains visibility as you scale. Monthly executive updates should highlight breakthrough applications,

adoption metrics trending in the right direction, and ROI calculations that justify continued investment. But also create internal channels—Slack channels, newsletters, all-hands segments—where employees see peers achieving results. Social proof drives adoption faster than top-down mandates.

Benchmark against industry standards

Partner with industry groups, consult analyst reports, or engage peer networks to understand where your AI maturity ranks compared to the rest of your market. Knowing you're outpacing competitors in adoption velocity or ROI per use case validates your strategy and identifies areas needing acceleration.

Establish feedback loops that capture qualitative insights alongside quantitative metrics. Survey users quarterly about what's working, what's frustrating, and what capabilities they wish existed. These insights reveal where to double down on successful patterns and which approaches need course correction before they scale to enterprise-wide problems.

Customer spotlight: [Asana](#), the work management platform, uses Claude to power AI features across their product. By systematically tracking user engagement with AI features and demonstrating how they accelerate project timelines, they've built both internal conviction and compelling customer case studies that drive product-led growth.



Chapter 4

Building an AI-first enterprise: real-world examples from Anthropic

Building an AI-first enterprise: real-world examples from Anthropic

To move beyond AI experimentation and into actually building AI-native operations, the most dramatic gains come not from isolated use cases, but from reimagining entire workflows around AI capabilities. At Anthropic, we've transformed how we work by integrating Claude throughout our company operations. Teams from Engineering to People Operations use Claude to build agents that work across long time horizons, manage in-depth knowledge bases, and tackle function-specific tasks using the same core agentic building blocks.

The examples below show what becomes possible when teams have the right AI tools and the autonomy to reimagine their processes.

Engineering

Anthropic's Engineering teams use Claude Code to dramatically accelerate software development cycles, eliminating the time-intensive work of building context in unfamiliar codebases. At Anthropic, engineers leverage Claude Code as their first stop for any development task, from bug fixes to feature development to running analysis.

When rotating onto new teams or working in unfamiliar code, engineers use Claude Code to quickly understand which files need modification and what changes are required, eliminating hours of manual navigation and the tedious process of copying snippets into Claude.ai for context. The impact extends beyond speed to fundamentally changing what engineers are willing to

tackle—developers now confidently dive into unfamiliar parts of the codebase to fix bugs or investigate incidents that they would have previously delegated because building context would take too long.

Revenue Operations

Anthropic's Revenue Operations team transformed work that typically takes months into days with Claude Code. One RevOps engineer built an entire Salesforce application in under a week—including an org chart solution that would typically cost hundreds of thousands annually—and created automated email systems that pull first-party usage data from Databricks to generate personalized account alerts with charts and insights. Claude Code enabled the rapid development of Slack notification systems when customers create new Anthropic organizations, ensuring sales teams stay informed in real-time.

Sales

Account executives use Claude for Work daily for comprehensive account management, maintaining project-based customer knowledge bases that include all meeting notes, preparation documents, and relevant research. What previously took six hours to prepare—executive briefing documents—now takes one hour with Claude's assistance. Claude excels at transforming rough drafts into polished, executive-ready communications while maintaining appropriate tone and messaging.

Legal

Rather than routing every question through legal teams, employees use Claude as a preliminary legal assistant to review terms of service, understand contract changes, and draft compliance responses for internal review. This allows legal teams to focus on complex negotiations while routine inquiries get immediate attention through AI assistance, significantly reducing ticket volume to legal departments.

Marketing

Marketing teams use Claude to accelerate content development across channels while maintaining brand consistency. Content creators draft blog posts, social media campaigns, and product announcements with Claude's assistance, completing in hours what previously required days of iteration. The team even built custom instructions that encode Anthropic's voice and style guidelines, ensuring all Claude-generated content aligns with brand standards. Campaign planning benefits from Claude's ability to generate multiple creative variations quickly, helping teams test different messaging approaches and identify what resonates with target audiences before committing resources to production.

Finance

Finance teams leverage Claude to transform data analysis and reporting workflows. Financial analysts use Claude to automate routine reporting tasks, generating executive summaries from complex datasets that previously required days of manual compilation. The team built custom prompts that extract key metrics from quarterly results, flag anomalies requiring investigation, and draft narrative explanations of variance analysis. Finance partners with business units more effectively by using Claude to translate technical financial concepts into clear business language.

Human Resources

Anthropic developed Ask Claude, an internal knowledge assistant that connects to all company documentation sources through MCP (Model Context Protocol) integrations. New hires use Ask Claude during onboarding to answer questions they might hesitate to ask colleagues, like "What does this acronym stand for?" or "How do I get an employee verification letter?" What previously required submitting tickets that took days to resolve now gets answered in minutes, allowing HR teams to focus on strategic initiatives rather than fielding repetitive questions.

The impact extends beyond just faster answers. Time to productivity for new hires improved significantly—employees can now become effective contributors within their first week by accessing institutional knowledge instantly. For Anthropic's lean people operations team, Ask Claude reduces ticket volume substantially, preventing burnout and allowing HR professionals to spend time on higher-value work like strategic planning and employee development. The system particularly benefits technical teams who can debug issues and find technical documentation without waiting on IT support queues.



Chapter 5

Your AI transformation starts today

Your AI transformation starts today

Successful transformation begins with an accurate assessment of your current capabilities and a clear understanding of the path forward

Organizational readiness assessment

Before embarking on your first few pilots, conduct an honest evaluation of where your organization stands. Use this matrix to assess readiness across critical dimensions and determine your optimal transformation path:

Dimension	Building foundation (1-2)	Growing capability (3-4)	Transformation ready (5-6)	Your Score
Executive commitment	AI viewed as IT project	CEO interested; competing priorities	CEO championing; multi-year commitment	--
Data infrastructure	Legacy systems; limited cloud	Hybrid cloud; basic DevOps	Cloud-native; strong engineering team	--
Technical capabilities	Initiatives frequently stall	Mixed track record; moderate adoption	Proven success; high trust culture	--
Change management	Departmental silos	Regular meetings; shared goals emerging	Integrated teams; aligned incentives	--
Cross-functional collaboration	No AI experience; exploratory	Initial ML models in production	Integrated teams; aligned incentives	--
AI/ML maturity	No AI experience; exploratory	Initial ML models in production	Multiple AI applications deployed	--
Risk & compliance	Reactive; manual controls	Established program; regular audits	Proactive; automated controls	--
Budget & resources	Project-based funding	Annual AI budget established	Multi-year investment secured	--

Scoring guide:

- **30-48 points (High Readiness):** Launch comprehensive transformation with multiple pilots across functions.
- **16-29 points (Moderate Readiness):** Begin with 3-5 strategic pilots while addressing foundational gaps.
- **8-15 points (Building Readiness):** Secure executive sponsorship and establish governance before launching 1-2 narrow pilots.

The vision: Make AI your competitive advantage

Six months from now, your enterprise could look radically different: Engineers shipping features 5x faster with AI pair programmers that understand your entire codebase. Customer service teams resolve complex issues in minutes instead of days, with AI agents that learn from every interaction. Financial analysts running sophisticated models that once required teams of quantitative experts.

But the real transformation goes deeper than task automation. AI-first organizations make better decisions faster, experiment continuously, and unlock insights buried in decades of institutional knowledge. They turn every employee into a power user and every process into an opportunity for innovation.

The question isn't whether AI will transform your industry—it's whether your company will lead that transformation.

Ready to begin your AI transformation?

Join enterprises like IG Group, Novo Nordisk, Cox Automotive, Palo Alto Networks, and Salesforce who are already transforming their operations with Claude.

Start building today:

Contact Anthropic's Sales team to discuss your specific use cases and implementation strategy

Access Claude through our API, web interface, or Claude for Work to begin piloting with your teams

Review our enterprise documentation and safety guidelines at docs.anthropic.com and our Trust Center

Explore Claude Code to accelerate engineering workflows



<https://claude.com/product/overview>